

# Exploring and pre-processing a dataset using Pandas

Estimated time needed: **30** minutes

## Objectives

After completing this lab you will be able to:

- Explore the dataset
- Pre-process dataset as required (may be for visualization)

## Introduction

The aim of this lab is to provide you a refresher on the **Pandas** library, so that you can pre-process and analyse the datasets before applying data visualization techniques on it. This lab will work as a crash course on *pandas*. If you are interested in learning more about the *pandas* library, detailed description and explanation of how to use it and how to clean, munge, and process data stored in a *pandas* dataframe are provided in other IBM courses.

## Table of Contents

1. Exploring Datasets with *pandas*
2. The Dataset: Immigration to Canada from 1980 to 2013
3. *pandas* Basics
4. *pandas* Intermediate: Indexing and Selection
5. *pandas* Filtering based on a criteria
6. *pandas* Sorting Values

## Exploring Datasets with *pandas*

*pandas* is an essential data analysis toolkit for Python. From their [website](#):

*pandas* is a Python package providing fast, flexible, and expressive data structures designed to make working with "relational" or "labeled" data both easy and intuitive. It aims to be the fundamental high-level building block for doing practical, **real world** data analysis in Python.

The course heavily relies on *pandas* for data wrangling, analysis, and visualization. We encourage you to spend some time and familiarize yourself with the *pandas* API. Reference: <http://pandas.pydata.org/pandas-docs/stable/api.html>.

## The Dataset: Immigration to Canada from 1980 to 2013

Dataset Source: [International migration flows to and from selected countries - The 2015 revision](#).

The dataset contains annual data on the flows of international immigrants as recorded by the countries of destination. The data presents both inflows and outflows according to the place of birth, citizenship or place of previous / next residence both for foreigners and nationals. The current version presents data pertaining to 45 countries.

In this lab, we will focus on the Canadian immigration data.

```
In [5]: df_can.tail()
```

```
Out[5]:
```

	Type	Coverage	OdName	AREA	AreaName	REG	RegName	DEV	DevName	1980	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
190	Immigrants	Foreigners	Viet Nam	935	Asia	920	South-Eastern Asia	902	Developing regions	1191	...	1816	1852	3153	2574	1784	2171	1942	1723	1731	2013
191	Immigrants	Foreigners	Western Sahara	903	Africa	912	Northern Africa	902	Developing regions	0	...	0	0	1	0	0	0	0	0	0	0
192	Immigrants	Foreigners	Yemen	935	Asia	922	Western Asia	902	Developing regions	1	...	124	161	140	122	133	128	211	160	174	2013
193	Immigrants	Foreigners	Zambia	903	Africa	910	Eastern Africa	902	Developing regions	11	...	56	91	77	71	64	60	102	69	46	2013
194	Immigrants	Foreigners	Zimbabwe	903	Africa	910	Eastern Africa	902	Developing regions	72	...	1450	615	454	663	611	508	494	434	437	2013

5 rows x 43 columns

When analyzing a dataset, it's always a good idea to start by getting basic information about your dataframe. We can do this by using the `info()` method.

This method can be used to get a short summary of the dataframe.

```
In [6]: df_can.info(verbose=False)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 195 entries, 0 to 194
Columns: 43 entries, Type to 2013
dtypes: int64(37), object(6)
memory usage: 65.6+ KB
```

To get the list of column headers we can call upon the dataframe's `columns` instance variable.

```
In [7]: df_can.columns
```

```
Out[7]: Index([      'Type',      'Coverage',      'OdName',      'AREA',      'AreaName',      'REG',
      'RegName',      'DEV',      'DevName',
      1980,      1981,      1982,
      1983,      1984,      1985,
      1986,      1987,      1988,
      1989,      1990,      1991,
      1992,      1993,      1994,
      1995,      1996,      1997,
      1998,      1999,      2000,
      2001,      2002,      2003,
      2004,      2005,      2006,
      2007,      2008,      2009,
      2010,      2011,      2012,
      2013],
              dtype='object')
```

Similarly, to get the list of indices we use the `index` instance variables.

```
In [8]: df_can.index
```

```
Out[8]: RangeIndex(start=0, stop=195, step=1)
```

Note: The default type of instance variables `index` and `columns` are **NOT** list.

```
In [9]: print(type(df_can.columns))
print(type(df_can.index))

<class 'pandas.core.indexes.base.Index'>
<class 'pandas.core.indexes.range.RangeIndex'>
```

To get the index and columns as lists, we can use the `tolist()` method.

```
In [ ]: df_can.columns.tolist()
```

```
In [ ]: df_can.index.tolist()
```

```
In [11]: print(type(df_can.columns.tolist()))
print(type(df_can.index.tolist()))

<class 'list'>
<class 'list'>
```

To view the dimensions of the dataframe, we use the `shape` instance variable of it.

```
In [12]: # size of dataframe (rows, columns)
df_can.shape
```

```
Out[12]: (195, 43)
```

**Note:** The main types stored in *pandas* objects are `float`, `int`, `bool`, `datetime64[ns]`, `datetime64[ns, tz]`, `timedelta[ns]`, `category`, and `object` (string). In addition, these dtypes have item sizes, e.g. `int64` and `int32`.

Let's clean the data set to remove a few unnecessary columns. We can use *pandas* `drop()` method as follows:

```
In [13]: # in pandas axis=0 represents rows (default) and axis=1 represents columns.
df_can.drop(['AREA', 'REG', 'DEV', 'Type', 'Coverage'], axis=1, inplace=True)
df_can.head(2)
```

```
Out[13]:
```

	OdName	AreaName	RegName	DevName	1980	1981	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
0	Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	...	2978	3436	3009	2652	2111	1746	1758	2203	2635	2004
1	Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	...	1450	1223	856	702	560	716	561	539	620	603

2 rows x 38 columns

Let's rename the columns so that they make sense. We can use `rename()` method by passing in a dictionary of old and new names as follows:

```
In [14]: df_can.rename(columns={'OdName': 'Country', 'AreaName': 'Continent', 'RegName': 'Region'}, inplace=True)
df_can.columns

Out[14]: Index([ 'Country', 'Continent', 'Region', 'DevName', 1980,
                1981, 1982, 1983, 1984, 1985,
                1986, 1987, 1988, 1989, 1990,
                1991, 1992, 1993, 1994, 1995,
                1996, 1997, 1998, 1999, 2000,
                2001, 2002, 2003, 2004, 2005,
                2006, 2007, 2008, 2009, 2010,
                2011, 2012, 2013],
              dtype='object')
```

We will also add a 'Total' column that sums up the total immigrants by country over the entire period 1980 - 2013, as follows:

```
In [15]: df_can['Total'] = df_can.sum(axis=1)
df_can['Total']

/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
  """Entry point for launching an IPython kernel.

Out[15]: 0      58639
         1      15699
         2      69439
         3         6
         4        15
         ...
        190     97146
        191         2
        192      2985
        193      1677
        194      8598
         Name: Total, Length: 195, dtype: int64
```

We can check to see how many null objects we have in the dataset as follows:

```
In [ ]: df_can.isnull().sum()
```

Finally, let's view a quick summary of each column in our dataframe using the `describe()` method.

```
In [17]: df_can.describe()

Out[17]:
```

	1980	1981	1982	1983	1984	1985	1986	1987	1988	1989	...	2005
count	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000	195.000000	...	195.000000
mean	508.394872	566.989744	534.723077	387.435897	376.497436	358.861538	441.271795	691.133333	714.389744	843.241026	...	1320.292308
std	1949.588546	2152.643752	1866.997511	1204.333597	1198.246371	1079.309600	1225.576630	2109.205607	2443.606788	2555.048874	...	4425.957828
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.500000	0.500000	1.000000	1.000000	...	28.500000
50%	13.000000	10.000000	11.000000	12.000000	13.000000	17.000000	18.000000	26.000000	34.000000	44.000000	...	210.000000
75%	251.500000	295.500000	275.000000	173.000000	181.000000	197.000000	254.000000	434.000000	409.000000	508.500000	...	832.000000
max	22045.000000	24796.000000	20620.000000	10015.000000	10170.000000	9564.000000	9470.000000	21337.000000	27359.000000	23795.000000	...	42584.000000

8 rows x 35 columns

# pandas Intermediate: Indexing and Selection (slicing)

## Select Column

There are two ways to filter on a column name:

Method 1: Quick and easy, but only works if the column name does NOT have spaces or special characters.

```
df.column_name # returns series
```

Method 2: More robust, and can filter on multiple columns.

```
df['column'] # returns series
df[['column 1', 'column 2']] # returns dataframe
```

Example: Let's try filtering on the list of countries ('Country').

```
In [ ]: df_can.Country # returns a series
```

Let's try filtering on the list of countries ('Country') and the data for years: 1980 - 1985.

```
In [19]: df_can[['Country', 1980, 1981, 1982, 1983, 1984, 1985]] # returns a dataframe
# notice that 'Country' is string, and the years are integers.
# for the sake of consistency, we will convert all column names to string later on.
```

```
Out[19]:
```

	Country	1980	1981	1982	1983	1984	1985
0	Afghanistan	16	39	39	47	71	340
1	Albania	1	0	0	0	0	0
2	Algeria	80	67	71	69	63	44
3	American Samoa	0	1	0	0	0	0
4	Andorra	0	0	0	0	0	0
...	...	...	...	...	...	...	...
190	Viet Nam	1191	1829	2162	3404	7583	5907
191	Western Sahara	0	0	0	0	0	0
192	Yemen	1	2	1	6	0	18
193	Zambia	11	17	11	7	16	9
194	Zimbabwe	72	114	102	44	32	29

195 rows × 7 columns

## Select Row

There are main 2 ways to select rows:

```
df.loc[label] # filters by the labels of the index/column
df.iloc[index] # filters by the positions of the index/column
```

Before we proceed, notice that the default index of the dataset is a numeric range from 0 to 194. This makes it very difficult to do a query by a specific country. For example to search for data on Japan, we need to know the corresponding index value.

This can be fixed very easily by setting the 'Country' column as the index using `set_index()` method.

```
In [20]: df_can.set_index('Country', inplace=True)
# tip: The opposite of set is reset. So to reset the index, we can use df_can.reset_index()
```

```
In [21]: df_can.head(3)
```

```
Out[21]:
```

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
Country																					
Afghanistan	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	496	...	3436	3009	2652	2111	1746	1758	2203	2635	2004	58639
Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	1	...	1223	856	702	560	716	561	539	620	603	15699
Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	69	...	3626	4807	3623	4005	5393	4752	4325	3774	4331	69439

3 rows × 38 columns

```
In [22]: # optional: to remove the name of the index
df_can.index.name = None
```

Example: Let's view the number of immigrants from Japan (row 87) for the following scenarios: 1. The full row data (all columns) 2. For year 2013 3. For years 1980 to 1985

```
In [24]: # 1. the full row data (all columns)
df_can.loc['Japan'][:5]
```

```
Out[24]:
```

Continent	Asia
Region	Eastern Asia
DevName	Developed regions
1980	701
1981	756
Name: Japan, dtype: object	

```
In [26]: # alternate methods
df_can.iloc[87][:5]
```

```
Out[26]: Continent      Asia
Region      Eastern Asia
DevName      Developed regions
1980          701
1981          756
Name: Japan, dtype: object
```

```
In [27]: df_can[df_can.index == 'Japan']
```

```
Out[27]:
```

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
Japan	Asia	Eastern Asia	Developed regions	701	756	598	309	246	198	248	...	1067	1212	1250	1284	1194	1168	1265	1214	982	27707

1 rows x 38 columns

```
In [28]: # 2. for year 2013
df_can.loc['Japan', 2013]
```

```
Out[28]: 982
```

```
In [29]: # alternate method
# year 2013 is the last column, with a positional index of 36
df_can.iloc[87, 36]
```

```
Out[29]: 982
```

```
In [30]: # 3. for years 1980 to 1985
df_can.loc['Japan', [1980, 1981, 1982, 1983, 1984, 1984]]
```

```
Out[30]: 1980    701
1981    756
1982    598
1983    309
1984    246
1984    246
Name: Japan, dtype: object
```

```
In [31]: # Alternative Method
df_can.iloc[87, [3, 4, 5, 6, 7, 8]]
```

```
Out[31]: 1980    701
1981    756
1982    598
1983    309
1984    246
1985    198
Name: Japan, dtype: object
```

**Exercise:** Let's view the number of immigrants from **Haiti** for the following scenarios:

1. The full row data (all columns)
2. For year 2000
3. For years 1990 to 1995

```
In [32]: df_can.loc['Haiti']
df_can.loc['Haiti', 2000]
df_can.loc['Haiti', [1990, 1991, 1992, 1993, 1994, 1995]]
```

```
Out[32]: 1990    2379
1991    2829
1992    2399
1993    3655
1994    2100
1995    2014
Name: Haiti, dtype: object
```

► [Click here for a sample python solution](#)

Column names that are integers (such as the years) might introduce some confusion. For example, when we are referencing the year 2013, one might confuse that when the 2013th positional index.

To avoid this ambiguity, let's convert the column names into strings: '1980' to '2013'.

```
In [33]: df_can.columns = list(map(str, df_can.columns))
# [print (type(x)) for x in df_can.columns.values] #<-- uncomment to check type of column headers
```

Since we converted the years to string, let's declare a variable that will allow us to easily call upon the full range of years:

```
In [35]: # useful for plotting later on
years = list(map(str, range(1980, 2014)))
print(years)
```

```
['1980', '1981', '1982', '1983', '1984', '1985', '1986', '1987', '1988', '1989', '1990', '1991', '1992', '1993', '1994', '1995', '1996', '1997', '1998', '1999', '2000', '2001', '2002', '2003', '2004', '2005', '2006', '2007', '2008', '2009', '2010', '2011', '2012', '2013']
```

**Exercise:** Create a list named 'year' using map function for years ranging from 1990 to 2013.

Then extract the data series from the dataframe df\_can for Haiti using year list.

```
In [38]: year = list(map(str, range(1990,2014)))
```

```
df_can[year].head()
```

Out[38]:

	1990	1991	1992	1993	1994	1995	1996	1997	1998	1999	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
<b>Afghanistan</b>	1028	1378	1170	713	858	1537	2212	2555	1999	2395	...	2978	3436	3009	2652	2111	1746	1758	2203	2635	2004
<b>Albania</b>	3	21	56	96	71	63	113	307	574	1264	...	1450	1223	856	702	560	716	561	539	620	603
<b>Algeria</b>	491	872	795	717	595	1106	2054	1842	2292	2389	...	3616	3626	4807	3623	4005	5393	4752	4325	3774	4331
<b>American Samoa</b>	2	0	0	0	0	0	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0
<b>Andorra</b>	3	0	1	0	0	0	0	0	2	0	...	0	0	1	1	0	0	0	0	1	1

5 rows × 24 columns

► [Click here for a sample python solution](#)

## Filtering based on a criteria

To filter the dataframe based on a condition, we simply pass the condition as a boolean vector.

For example, Let's filter the dataframe to show the data on Asian countries (AreaName = Asia).

```
In [39]: # 1. create the condition boolean series
condition = df_can['Continent'] == 'Asia'
print(condition)
```

```
Afghanistan      True
Albania           False
Algeria           False
American Samoa   False
Andorra          False
...
Viet Nam         True
Western Sahara   False
Yemen            True
Zambia           False
Zimbabwe         False
Name: Continent, Length: 195, dtype: bool
```

```
In [41]: # 2. pass this condition into the dataframe
df_can[condition][:5]
```

Out[41]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
<b>Afghanistan</b>	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	496	...	3436	3009	2652	2111	1746	1758	2203	2635	2004	58639
<b>Armenia</b>	Asia	Western Asia	Developing regions	0	0	0	0	0	0	0	...	224	218	198	205	267	252	236	258	207	3310
<b>Azerbaijan</b>	Asia	Western Asia	Developing regions	0	0	0	0	0	0	0	...	359	236	203	125	165	209	138	161	57	2649
<b>Bahrain</b>	Asia	Western Asia	Developing regions	0	2	1	1	1	3	0	...	12	12	22	9	35	28	21	39	32	475
<b>Bangladesh</b>	Asia	Southern Asia	Developing regions	83	84	86	81	98	92	486	...	4171	4014	2897	2939	2104	4721	2694	2640	3789	65568

5 rows × 38 columns

```
In [42]: # we can pass multiple criteria in the same line.
# Let's filter for AreaName = Asia and RegName = Southern Asia

df_can[(df_can['Continent']=='Asia') & (df_can['Region']=='Southern Asia')]

# note: When using 'and' and 'or' operators, pandas requires we use '&' and '|' instead of 'and' and 'or'
# don't forget to enclose the two conditions in parentheses
```

Out[42]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	
<b>Afghanistan</b>	Asia	Southern Asia	Developing regions	16	39	39	47	71	340	496	...	3436	3009	2652	2111	1746	1758	2203	2635	2004	5
<b>Bangladesh</b>	Asia	Southern Asia	Developing regions	83	84	86	81	98	92	486	...	4171	4014	2897	2939	2104	4721	2694	2640	3789	6
<b>Bhutan</b>	Asia	Southern Asia	Developing regions	0	0	0	0	1	0	0	...	5	10	7	36	865	1464	1879	1075	487	
<b>India</b>	Asia	Southern Asia	Developing regions	8880	8670	8147	7338	5704	4211	7150	...	36210	33848	28742	28261	29456	34235	27509	30933	33087	69
<b>Iran (Islamic Republic of)</b>	Asia	Southern Asia	Developing regions	1172	1429	1822	1592	1977	1648	1794	...	5837	7480	6974	6475	6580	7477	7479	7534	11291	17
<b>Maldives</b>	Asia	Southern Asia	Developing regions	0	0	0	1	0	0	0	...	0	0	2	1	7	4	3	1	1	
<b>Nepal</b>	Asia	Southern Asia	Developing regions	1	1	6	1	2	4	13	...	607	540	511	581	561	1392	1129	1185	1308	1
<b>Pakistan</b>	Asia	Southern Asia	Developing regions	978	972	1201	900	668	514	691	...	14314	13127	10124	8994	7217	6811	7468	11227	12603	24
<b>Sri Lanka</b>	Asia	Southern Asia	Developing regions	185	371	290	197	1086	845	1838	...	4930	4714	4123	4756	4547	4422	3309	3338	2394	14

9 rows × 38 columns

**Exercise:** Fetch the data where AreaName is 'Africa' and RegName is 'Southern Africa'.  
Display the dataframe and find out how many instances are there?

In [44]:

df\_can[(df\_can['Continent']=='Africa') & (df\_can['Region']=='Southern Africa')][:5]

Out[44]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
<b>Botswana</b>	Africa	Southern Africa	Developing regions	10	1	3	3	7	4	2	...	7	11	8	28	15	42	53	64	76	396
<b>Lesotho</b>	Africa	Southern Africa	Developing regions	1	1	1	2	7	5	3	...	4	0	4	1	8	7	1	0	6	107
<b>Namibia</b>	Africa	Southern Africa	Developing regions	0	5	5	3	2	1	1	...	6	19	13	26	14	16	23	24	83	320
<b>South Africa</b>	Africa	Southern Africa	Developing regions	1026	1118	781	379	271	310	718	...	988	1111	1200	1123	1188	1238	959	1243	1240	40568
<b>Swaziland</b>	Africa	Southern Africa	Developing regions	4	1	1	0	10	7	1	...	7	7	5	6	10	3	13	17	39	188

5 rows × 38 columns

► [Click here for a sample python solution](#)

### Sorting Values of a Dataframe or Series

You can use the `sort_values()` function is used to sort a DataFrame or a Series based on one or more columns.  
You to specify the column(s) by which you want to sort and the order (ascending or descending). Below is the syntax to use it:-

```
df.sort_values(col_name, axis=0, ascending=True, inplace=False, ignore_index=False)
```

- col\_nam - the column(s) to sort by.
- axis - axis along which to sort. 0 for sorting by rows (default) and 1 for sorting by columns.
- ascending - to sort in ascending order (True, default) or descending order (False).
- inplace - to perform the sorting operation in-place (True) or return a sorted copy (False, default).
- ignore\_index - to reset the index after sorting (True) or keep the original index values (False, default).

Let's sort out dataframe df\_can on 'Total' column, in descending order to find out the top 5 countries that contributed the most to immigration to Canada.

In [45]:

df\_can.sort\_values(by='Total', ascending=False, axis=0, inplace=True)  
top\_5 = df\_can.head(5)  
top\_5



Out[45]:

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	
	India	Asia	Southern Asia	Developing regions	8880	8670	8147	7338	5704	4211	7150	...	36210	33848	28742	28261	29456	34235	27509	30933	33087
	China	Asia	Eastern Asia	Developing regions	5123	6682	3308	1863	1527	1816	1960	...	42584	33518	27642	30037	29622	30391	28502	33024	34129
	United Kingdom of Great Britain and Northern Ireland	Europe	Northern Europe	Developed regions	22045	24796	20620	10015	10170	9564	9470	...	7258	7140	8216	8979	8876	8724	6204	6195	5827
	Philippines	Asia	South-Eastern Asia	Developing regions	6051	5921	5249	4562	3801	3150	4166	...	18139	18400	19837	24887	28573	38617	36765	34315	29544
	Pakistan	Asia	Southern Asia	Developing regions	978	972	1201	900	668	514	691	...	14314	13127	10124	8994	7217	6811	7468	11227	12603

5 rows × 38 columns

**Exercise:** Find out top 3 countries that contributes the most to immigration to Canda in the year 2010.  
Display the country names with the immigrant count in this year

In [47]:

```
df_can.sort_values(by='2010', ascending=False, axis=0, inplace=True)
top_3 = df_can['2010'].head(3)
top_3
```

Out[47]:

```
Philippines    38617
India          34235
China          30391
Name: 2010, dtype: int64
```

► Click here for a sample python solution

Congratulations! you have learned how to wrangle data with Pandas. You will be using alot of these commands to preprocess the data before its can be used for data visualization.

Thank you for completing this lab!

Author

Alex Aklson

Other Contributors

Jay Rajasekharan, Ehsan M. Kermani, Slobodan Markovic, Weiqing Wang, Dr. Pooja

<!-- --!> ## Change Log | Date (YYYY-MM-DD) | Version | Changed By | Change Description |  
|-----|-----|-----|-----| 2023-06-08 | 2.5 | Dr. Pooja | Separated from original lab | | 2021-05-29 | 2.4 | Weiqing Wang | Fixed typos and code smells. | | 2021-01-20 | 2.3 | Lakshmi Holla | Changed TOC cell markdown | | 2020-11-20 | 2.2 | Lakshmi Holla | Changed IBM box URL | | 2020-11-03 | 2.1 | Lakshmi Holla | Changed URL and info method | | 2020-08-27 | 2.0 | Lavanya | Moved Lab to course repo in GitLab |--!> ##